

scipy / scipy

Branch: master ▾

scipy / scipy / optimize / lbfgsb.py

Find file

Copy path

pv ENH: optimize: call callback with a copy for scipy.optimize

7f43e34 4 days ago

19 contributors



469 lines (397 sloc) 16.5 KB

```

1  """
2  Functions
3  -----
4  .. autosummary::
5      :toctree: generated/
6
7      fmin_l_bfgs_b
8
9  """
10
11  ## License for the Python wrapper
12  ## =====
13
14  ## Copyright (c) 2004 David M. Cooke <cookedm@physics.mcmaster.ca>
15
16  ## Permission is hereby granted, free of charge, to any person obtaining a
17  ## copy of this software and associated documentation files (the "Software"),
18  ## to deal in the Software without restriction, including without limitation
19  ## the rights to use, copy, modify, merge, publish, distribute, sublicense,
20  ## and/or sell copies of the Software, and to permit persons to whom the
21  ## Software is furnished to do so, subject to the following conditions:
22
23  ## The above copyright notice and this permission notice shall be included in
24  ## all copies or substantial portions of the Software.
25
26  ## THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR
27  ## IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY,
28  ## FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE
29  ## AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER
30  ## LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING
31  ## FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER
32  ## DEALINGS IN THE SOFTWARE.
33
34  ## Modifications by Travis Oliphant and Enthought, Inc. for inclusion in SciPy
35
36  from __future__ import division, print_function, absolute_import
37
38  import numpy as np
39  from numpy import array, asarray, float64, int32, zeros
40  from . import lbfgsb
41  from .optimize import (approx_fprime, MemoizeJac, OptimizeResult,
42                        _check_unknown_options, wrap_function,
43                        _approx_fprime_helper)
44  from scipy.sparse.linalg import LinearOperator
45
46  __all__ = ['fmin_l_bfgs_b', 'LbfgsInvHessProduct']
47
48
49  def fmin_l_bfgs_b(func, x0, fprime=None, args=(),
50                  approx_grad=0,
51                  bounds=None, m=10, factr=1e7, pgtol=1e-5,
52                  epsilon=1e-8,
53                  iprint=-1, maxfun=15000, maxiter=15000, disp=None,
54                  callback=None, maxls=20):
55
56      """
57      Minimize a function func using the L-BFGS-B algorithm.
58
59      Parameters
60      -----
61      func : callable f(x,*args)
62          Function to minimise.
63      x0 : ndarray

```

```

63     Initial guess.
64 fprime : callable fprime(x,*args), optional
65     The gradient of `func`. If None, then `func` returns the function
66     value and the gradient (`f, g = func(x, *args)`), unless
67     `approx_grad` is True in which case `func` returns only `f`.
68 args : sequence, optional
69     Arguments to pass to `func` and `fprime`.
70 approx_grad : bool, optional
71     Whether to approximate the gradient numerically (in which case
72     `func` returns only the function value).
73 bounds : list, optional
74     ``(min, max)`` pairs for each element in ``x``, defining
75     the bounds on that parameter. Use None or +-inf for one of ``min`` or
76     ``max`` when there is no bound in that direction.
77 m : int, optional
78     The maximum number of variable metric corrections
79     used to define the limited memory matrix. (The limited memory BFGS
80     method does not store the full hessian but uses this many terms in an
81     approximation to it.)
82 factr : float, optional
83     The iteration stops when
84     ``(f^k - f^{k+1})/max{|f^k|, |f^{k+1}|, 1} <= factr * eps``,
85     where ``eps`` is the machine precision, which is automatically
86     generated by the code. Typical values for `factr` are: 1e12 for
87     low accuracy; 1e7 for moderate accuracy; 10.0 for extremely
88     high accuracy. See Notes for relationship to `ftol`, which is exposed
89     (instead of `factr`) by the `scipy.optimize.minimize` interface to
90     L-BFGS-B.
91 pgtol : float, optional
92     The iteration will stop when
93     ``max{|proj g_i | i = 1, ..., n} <= pgtol``
94     where ``pg_i`` is the i-th component of the projected gradient.
95 epsilon : float, optional
96     Step size used when `approx_grad` is True, for numerically
97     calculating the gradient
98 iprint : int, optional
99     Controls the frequency of output. ``iprint < 0`` means no output;
100    ``iprint = 0`` print only one line at the last iteration;
101    ``0 < iprint < 99`` print also f and ``|proj g|`` every iprint iterations;
102    ``iprint = 99`` print details of every iteration except n-vectors;
103    ``iprint = 100`` print also the changes of active set and final x;
104    ``iprint > 100`` print details of every iteration including x and g.
105 disp : int, optional
106     If zero, then no output. If a positive number, then this over-rides
107     `iprint` (i.e., `iprint` gets the value of `disp`).
108 maxfun : int, optional
109     Maximum number of function evaluations.
110 maxiter : int, optional
111     Maximum number of iterations.
112 callback : callable, optional
113     Called after each iteration, as ``callback(xk)``, where ``xk`` is the
114     current parameter vector.
115 maxls : int, optional
116     Maximum number of line search steps (per iteration). Default is 20.
117
118 Returns
119 -----
120 x : array_like
121     Estimated position of the minimum.
122 f : float
123     Value of `func` at the minimum.
124 d : dict
125     Information dictionary.
126
127     * d['warnflag'] is
128
129         - 0 if converged,
130         - 1 if too many function evaluations or too many iterations,
131         - 2 if stopped for another reason, given in d['task']
132
133     * d['grad'] is the gradient at the minimum (should be 0 ish)
134     * d['funcalls'] is the number of function calls made.
135     * d['nit'] is the number of iterations.
136

```

```

137 See also
138 -----
139 minimize: Interface to minimization algorithms for multivariate
140 functions. See the 'L-BFGS-B' `method` in particular. Note that the
141 `ftol` option is made available via that interface, while `factr` is
142 provided via this interface, where `factr` is the factor multiplying
143 the default machine floating-point precision to arrive at `ftol`:
144 ``ftol = factr * numpy.finfo(float).eps``.
145
146 Notes
147 -----
148 License of L-BFGS-B (FORTRAN code):
149
150 The version included here (in fortran code) is 3.0
151 (released April 25, 2011). It was written by Ciyou Zhu, Richard Byrd,
152 and Jorge Nocedal <nocedal@ece.nwu.edu>. It carries the following
153 condition for use:
154
155 This software is freely available, but we expect that all publications
156 describing work using this software, or all commercial products using it,
157 quote at least one of the references given below. This software is released
158 under the BSD License.
159
160 References
161 -----
162 * R. H. Byrd, P. Lu and J. Nocedal. A Limited Memory Algorithm for Bound
163   Constrained Optimization, (1995), SIAM Journal on Scientific and
164   Statistical Computing, 16, 5, pp. 1190-1208.
165 * C. Zhu, R. H. Byrd and J. Nocedal. L-BFGS-B: Algorithm 778: L-BFGS-B,
166   FORTRAN routines for large scale bound constrained optimization (1997),
167   ACM Transactions on Mathematical Software, 23, 4, pp. 550 - 560.
168 * J.L. Morales and J. Nocedal. L-BFGS-B: Remark on Algorithm 778: L-BFGS-B,
169   FORTRAN routines for large scale bound constrained optimization (2011),
170   ACM Transactions on Mathematical Software, 38, 1.
171
172 """
173 # handle fprime/approx_grad
174 if approx_grad:
175     fun = func
176     jac = None
177 elif fprime is None:
178     fun = MemoizeJac(func)
179     jac = fun.derivative
180 else:
181     fun = func
182     jac = fprime
183
184 # build options
185 if disp is None:
186     disp = iprint
187 opts = {'disp': disp,
188        'iprint': iprint,
189        'maxcon': m,
190        'ftol': factr * np.finfo(float).eps,
191        'gtol': pgtol,
192        'eps': epsilon,
193        'maxfun': maxfun,
194        'maxiter': maxiter,
195        'callback': callback,
196        'maxls': maxls}
197
198 res = _minimize_lbfgsb(fun, x0, args=args, jac=jac, bounds=bounds,
199                      **opts)
200 d = {'grad': res['jac'],
201     'task': res['message'],
202     'funcalls': res['nfev'],
203     'nit': res['nit'],
204     'warnflag': res['status']}
205 f = res['fun']
206 x = res['x']
207
208 return x, f, d
209
210

```

```

211 def _minimize_lbfgsb(fun, x0, args=(), jac=None, bounds=None,
212                     disp=None, maxcor=10, ftol=2.2204460492503131e-09,
213                     gtol=1e-5, eps=1e-8, maxfun=15000, maxiter=15000,
214                     iprint=-1, callback=None, maxls=20, **unknown_options):
215     """
216     Minimize a scalar function of one or more variables using the L-BFGS-B
217     algorithm.
218
219     Options
220     -----
221     disp : bool
222         Set to True to print convergence messages.
223     maxcor : int
224         The maximum number of variable metric corrections used to
225         define the limited memory matrix. (The limited memory BFGS
226         method does not store the full hessian but uses this many terms
227         in an approximation to it.)
228     ftol : float
229         The iteration stops when ``|f^k -
230         f^{k+1}|/max{|f^k|, |f^{k+1}|, 1} <= ftol``.
231     gtol : float
232         The iteration will stop when ``max{|proj g_i | i = 1, ..., n}
233         <= gtol`` where ``pg_i`` is the i-th component of the
234         projected gradient.
235     eps : float
236         Step size used for numerical approximation of the jacobian.
237     disp : int
238         Set to True to print convergence messages.
239     maxfun : int
240         Maximum number of function evaluations.
241     maxiter : int
242         Maximum number of iterations.
243     maxls : int, optional
244         Maximum number of line search steps (per iteration). Default is 20.
245
246     Notes
247     -----
248     The option `ftol` is exposed via the `scipy.optimize.minimize` interface,
249     but calling `scipy.optimize.fmin_l_bfgs_b` directly exposes `factr`. The
250     relationship between the two is ``ftol = factr * numpy.finfo(float).eps``.
251     I.e., `factr` multiplies the default machine floating-point precision to
252     arrive at `ftol`.
253
254     """
255     _check_unknown_options(unknown_options)
256     m = maxcor
257     epsilon = eps
258     pgtol = gtol
259     factr = ftol / np.finfo(float).eps
260
261     x0 = asarray(x0).ravel()
262     n, = x0.shape
263
264     if bounds is None:
265         bounds = [(None, None)] * n
266     if len(bounds) != n:
267         raise ValueError('length of x0 != length of bounds')
268     # unbounded variables must use None, not +/-inf, for optimizer to work properly
269     bounds = [(None if l == -np.inf else l, None if u == np.inf else u) for l, u in bounds]
270
271     if disp is not None:
272         if disp == 0:
273             iprint = -1
274         else:
275             iprint = disp
276
277     n_function_evals, fun = wrap_function(fun, ())
278     if jac is None:
279         def func_and_grad(x):
280             f = fun(x, *args)
281             g = _approx_fprime_helper(x, fun, epsilon, args=args, f0=f)
282             return f, g
283     else:
284         def func_and_grad(x):

```

```

285         f = fun(x, *args)
286         g = jac(x, *args)
287         return f, g
288
289     nbd = zeros(n, int32)
290     low_bnd = zeros(n, float64)
291     upper_bnd = zeros(n, float64)
292     bounds_map = {(None, None): 0,
293                   (1, None): 1,
294                   (1, 1): 2,
295                   (None, 1): 3}
296     for i in range(0, n):
297         l, u = bounds[i]
298         if l is not None:
299             low_bnd[i] = l
300             l = 1
301         if u is not None:
302             upper_bnd[i] = u
303             u = 1
304         nbd[i] = bounds_map[l, u]
305
306     if not maxls > 0:
307         raise ValueError('maxls must be positive.')
308
309     x = array(x0, float64)
310     f = array(0.0, float64)
311     g = zeros((n,), float64)
312     wa = zeros(2*m*n + 5*n + 11*m*m + 8*m, float64)
313     iwa = zeros(3*n, int32)
314     task = zeros(1, 'S60')
315     csave = zeros(1, 'S60')
316     lsave = zeros(4, int32)
317     isave = zeros(44, int32)
318     dsave = zeros(29, float64)
319
320     task[:] = 'START'
321
322     n_iterations = 0
323
324     while 1:
325         # x, f, g, wa, iwa, task, csave, lsave, isave, dsave = \
326         _lbfgsb.setulb(m, x, low_bnd, upper_bnd, nbd, f, g, factr,
327                       pgtol, wa, iwa, task, iprint, csave, lsave,
328                       isave, dsave, maxls)
329         task_str = task.tostring()
330         if task_str.startswith(b'FG'):
331             # The minimization routine wants f and g at the current x.
332             # Note that interruptions due to maxfun are postponed
333             # until the completion of the current minimization iteration.
334             # Overwrite f and g:
335             f, g = func_and_grad(x)
336         elif task_str.startswith(b'NEW_X'):
337             # new iteration
338             n_iterations += 1
339             if callback is not None:
340                 callback(np.copy(x))
341
342             if n_iterations >= maxiter:
343                 task[:] = 'STOP: TOTAL NO. of ITERATIONS REACHED LIMIT'
344             elif n_function_evals[0] > maxfun:
345                 task[:] = ('STOP: TOTAL NO. of f AND g EVALUATIONS '
346                           'EXCEEDS LIMIT')
347             else:
348                 break
349
350     task_str = task.tostring().strip(b'\x00').strip()
351     if task_str.startswith(b'CONV'):
352         warnflag = 0
353     elif n_function_evals[0] > maxfun or n_iterations >= maxiter:
354         warnflag = 1
355     else:
356         warnflag = 2
357
358     # These two portions of the workspace are described in the mainlb

```

```

359 # subroutine in lbfgsb.f. See line 363.
360 s = wa[0: m*n].reshape(m, n)
361 y = wa[m*n: 2*m*n].reshape(m, n)
362
363 # See lbfgsb.f line 160 for this portion of the workspace.
364 # isave(31) = the total number of BFGS updates prior the current iteration;
365 n_bfgs_updates = isave[30]
366
367 n_corrs = min(n_bfgs_updates, maxcor)
368 hess_inv = LbfgsInvHessProduct(s[:n_corrs], y[:n_corrs])
369
370 return OptimizeResult(fun=f, jac=g, nfev=n_function_evals[0],
371                       nit=n_iterations, status=warnflag, message=task_str,
372                       x=x, success=(warnflag == 0), hess_inv=hess_inv)
373
374
375 class LbfgsInvHessProduct(LinearOperator):
376     """Linear operator for the L-BFGS approximate inverse Hessian.
377
378     This operator computes the product of a vector with the approximate inverse
379     of the Hessian of the objective function, using the L-BFGS limited
380     memory approximation to the inverse Hessian, accumulated during the
381     optimization.
382
383     Objects of this class implement the ``scipy.sparse.linalg.LinearOperator``
384     interface.
385
386     Parameters
387     -----
388     sk : array_like, shape=(n_corr, n)
389         Array of `n_corr` most recent updates to the solution vector.
390         (See [1]).
391     yk : array_like, shape=(n_corr, n)
392         Array of `n_corr` most recent updates to the gradient. (See [1]).
393
394     References
395     -----
396     .. [1] Nocedal, Jorge. "Updating quasi-Newton matrices with limited
397         storage." Mathematics of computation 35.151 (1980): 773-782.
398
399     """
400     def __init__(self, sk, yk):
401         """Construct the operator."""
402         if sk.shape != yk.shape or sk.ndim != 2:
403             raise ValueError('sk and yk must have matching shape, (n_corrs, n)')
404         n_corrs, n = sk.shape
405
406         super(LbfgsInvHessProduct, self).__init__(
407             dtype=np.float64, shape=(n, n))
408
409         self.sk = sk
410         self.yk = yk
411         self.n_corrs = n_corrs
412         self.rho = 1 / np.einsum('ij,ij->i', sk, yk)
413
414     def _matvec(self, x):
415         """Efficient matrix-vector multiply with the BFGS matrices.
416
417         This calculation is described in Section (4) of [1].
418
419         Parameters
420         -----
421         x : ndarray
422             An array with shape (n,) or (n,1).
423
424         Returns
425         -----
426         y : ndarray
427             The matrix-vector product
428
429         """
430         s, y, n_corrs, rho = self.sk, self.yk, self.n_corrs, self.rho
431         q = np.array(x, dtype=self.dtype, copy=True)
432         if q.ndim == 2 and q.shape[1] == 1:

```

```

433         q = q.reshape(-1)
434
435     alpha = np.zeros(n_corrs)
436
437     for i in range(n_corrs-1, -1, -1):
438         alpha[i] = rho[i] * np.dot(s[i], q)
439         q = q - alpha[i]*y[i]
440
441     r = q
442     for i in range(n_corrs):
443         beta = rho[i] * np.dot(y[i], r)
444         r = r + s[i] * (alpha[i] - beta)
445
446     return r
447
448 def todense(self):
449     """Return a dense array representation of this operator.
450
451     Returns
452     -----
453     arr : ndarray, shape=(n, n)
454         An array with the same shape and containing
455         the same data represented by this `LinearOperator`.
456
457     """
458     s, y, n_corrs, rho = self.sk, self.yk, self.n_corrs, self.rho
459     I = np.eye(*self.shape, dtype=self.dtype)
460     Hk = I
461
462     for i in range(n_corrs):
463         A1 = I - s[i][:, np.newaxis] * y[i][np.newaxis, :] * rho[i]
464         A2 = I - y[i][:, np.newaxis] * s[i][np.newaxis, :] * rho[i]
465
466         Hk = np.dot(A1, np.dot(Hk, A2)) + (rho[i] * s[i][:, np.newaxis] *
467                                           s[i][np.newaxis, :])
468     return Hk

```